

Loop Programming with /loop

This document describes the **/loop goal-driven loop programming extension** (spec: specs/agentloop). /loop lets you *program* the agent's behaviour with a declarative specification: a goal (success state), an optional verification command, a tool surface, scope boundaries, and budget limits. The agent then keeps its Reason -> Act -> Observe loop running until the specification says it is done.

/loop is **not** the core per-turn agent loop. The core loop is always running for every turn (see docs/howtos/reactagent.md). /loop *wraps* that loop: it injects stop conditions, a verification gate, workspace capture/rollback, and a "Goal Loop" section into the system prompt at fixed safe points.

Contents

1. What /loop adds
 2. Architecture overview
 3. Starting a loop
 4. The interactive setup dialog
 5. Writing a goal
 6. The LoopSpec fields
 7. Stop conditions and termination statuses
 8. The verification gate
 9. Restrictions: tool set, read-only, scope
 10. Checkpoints and rollback
 11. Configuration options
 12. Commands and keyboard surface
 13. Worked examples
 14. Events and telemetry
 15. Related documentation
-

1. What /loop adds

Without /loop, a turn works like this: you send a message, the model streams a response, any requested tools execute, results are appended, and the loop ends when the model answers without tool calls. You then prompt again for the next step.

With /loop, you declare an end state up front and the agent iterates autonomously toward it:

- **Goal** – the success state, expressed in plain text. Required.
- **Verification command** (optional) – a shell command whose *exit status* decides whether the goal is actually achieved. This is the strongest stop condition: a text-only "I'm done" response does not count until the command exits 0.
- **Tool set** (optional) – restricts which tools the loop may use.
- **Read-only constraints** (optional) – glob patterns for paths the loop must not write.
- **Scope boundaries** (optional) – glob patterns limiting where the loop may act at all.
- **Budgets** – a maximum iteration count and/or a maximum accumulated token cost.

Everything is enforced inside the normal agent loop at safe points, so a running loop remains interruptible and observable like any other turn.

2. Architecture overview

Layer	Location	Contents
State + enforcement	<code>crates/ragent-agent/src/session/loop_spec.rs</code>	<code>LoopSpec</code> , <code>stop_tracker</code> , <code>StopCondition</code> , <code>LOOP_WRITE_TOOLS</code> , <code>LOOP_ALWAYS_ALLOWED_TOOLS</code> , the <code>deny_reason</code> predicate
Runtime wiring	<code>crates/ragent-agent/src/session/processors.rs</code>	<code>SessionProcessors::start_loop</code> , <code>terminate_loop_inner</code> , <code>request_loop_interrupt</code> , <code>ensure_pre_loop_capture</code> , <code>rollback_loop</code> , <code>clear_loop</code> , <code>clear_loop_captures</code>
Verification	<code>crates/ragent-agent/src/session/verification.rs</code>	<code>VerificationCommand</code> , <code>VerificationOutcome</code> , <code>verification_failure_observation</code>
Capture / rollback	<code>crates/ragent-agent/src/session/capture.rs</code>	<code>PreLoopSnapshot</code> , <code>change summary</code> , <code>compute_change_summary</code>
Prompt injection	<code>crates/ragent-agent/src/session/prompt_injection.rs</code>	<code>GoalLoopSection</code> renders the <code>## Goal Loop</code> block into the turn system prompt
One-shot parsing	<code>crates/ragent-tui/src/app/shell.rs</code>	<code>handles_loop_command</code> (<code>/loop <agent> <goal...></code>)
Interactive dialog	<code>crates/ragent-tui/src/app/loop_setup.rs</code> (+ <code>loop_dialog/</code> module dir)	<code>LoopSetupField</code> , <code>LoopSetupState</code> , <code>build_spec_from_state</code> , <code>key handling</code>
Termination rendering	<code>crates/ragent-tui/src/app/helpers.rs</code> <code>event_handler.rs</code>	<code>loop_termination_banner</code> , <code>LoopTerminated / LoopChangeSummary handling</code> , <code>rollback offer</code>
Config	<code>crates/ragent-config/src/config.rs</code>	<code>LoopConfig</code> serde section (loop key) and its defaults

`SessionProcessor::start_loop` accepts an already-built `LoopSpec`; argument parsing and validation live in the TUI layer.

3. Starting a loop

There are three ways to start a goal loop.

3.1 One-shot (slash command)

```
/loop <agent> <goal text...>
```

Everything up to the first whitespace is the agent preset; the remainder is the goal. Examples:

```
/loop coder make all unit tests pass
```

```
/loop general summarise every README under docs/ into one index
```

One-shot defaults (documented in the handler doc comment):

- Step limit: the `loop.max_steps` value from `ragent.json` (default **512**).
- Cost limit: the `loop.cost_limit` value from `ragent.json` (default: none).
- Checkpoints: **on**.
- No verification command, no scope, no read-only constraints, no tool-set restriction.

Three one-shot flags override the run's budgets (any position, `--flag value` or `--flag=value`; `--cost_limit` and `--cost-limit` are both accepted):

```
/loop coder --max-steps 10 make all unit tests pass
```

```
/loop coder "refactor the auth module" --max-steps=40 --cost_limit 200000
```

```
/loop general --timeout 30 ship the release notes
```

Flag	Effect
<code>--max-steps N</code>	Override the step budget for this run
<code>--cost_limit N</code> / <code>--cost-limit N</code>	Override the token-cost budget for this run
<code>--timeout N</code>	Override the checkpoint-prompt timeout (seconds) for this run

An invalid flag value (missing, non-numeric, or out of range) stops the loop from starting and reports the offending flag. All other richer knobs (verify, scope, read-only, tool set) are set in the interactive dialog.

Starting a loop with an empty goal shows an error naming the missing field and re-opens the setup dialog with the typed agent pre-selected:

```
From: /loop
```

```
**Error:** missing field: goal - a loop cannot start without a goal.
```

```
Usage: `/loop <agent> <goal text...>` or `/loop` for the setup dialog.
```

```
Re-opening the setup dialog with your agent ({agent}) pre-selected.
```

The status bar shows `loop: missing field: goal`.

3.2 Interactive dialog

```
/loop
```

Opens the `/loop` setup overlay where all nine fields can be set (Section 4). Pressing Enter validates and starts the loop; Esc cancels but keeps every entered value as a draft for the next open.

3.3 Preconditions

App::start_goal_loop rejects the start (status-bar messages, no loop started) when:

- No provider is configured: loop: no provider configured - use /provider first
- No model is selected: loop: no model selected - use /model first
- No active session: [warn] No active session

On success the message window shows [loop {agent}] {goal}, the working status becomes loop running, and one spawned task runs processor.start_loop(&sid, spec) followed by processor.process_message(&sid, &goal_text, &agent, flag).

The **driver resolution** rule: the spec's agent preset wins; an unknown preset name falls back to the currently selected agent, so a typo never silently changes the driver.

4. The interactive setup dialog

4.1 Fields (in Tab order)

#	Field	Meaning	Notes
1	agent	Agent preset that drives the loop	Arrow keys select; selected agent is bold cyan
2	goal	Success state the loop must reach	Required (FR-005)
3	verify cmd	Verification command gating achievement	Optional (FR-007)
4	scope	Scope boundaries, comma-separated globs	Optional (FR-022)
5	read-only	Read-only constraints, comma-separated globs	Optional (FR-021)
6	tools	Restricted tool surface, comma-separated names	Optional (FR-008)
7	max steps	Step budget	Blank applies the config default (FR-013)
8	cost limit	Token cost budget	Blank applies the config default (FR-014)
9	checkpoints	Destructive-action checkpoint toggle	Toggled with Space or Up/Down; defaults on

max steps is prefilled from the on-disk loop.max_steps config value; cost limit is prefilled only when a value is configured.

4.2 Key bindings

Key	Action
Tab / BackTab	Next / previous field
Up / Down	agent: pick previous/next agent; checkpoints: toggle; text fields: move field
Space	Toggle checkpoints when focused
Enter	Validate and start the loop (draft slot cleared on success)
Esc	Close without starting; all entered values are kept as a draft (loop: setup cancelled (values kept))
Ctrl+V / Ctrl+W / Ctrl+K / Ctrl+U	Paste / delete previous word / delete to end / clear line (in text fields)
Shift+Arrow	Select text (in text fields)

If validation fails (missing goal), the error is rendered inside the dialog in red: `error: missing field: goal` - a loop cannot start without a goal.

The footer hint line reads: Tab next field - Up/Down navigate - Enter start - Esc cancel.

5. Writing a goal

A goal is a *success state*, not a task description. Phrase it so that a software process (or a human reviewer) can tell unambiguously whether the state has been reached. The strongest form pairs the goal with a verification command:

Element	Weak	Strong
Goal	“improve the tests”	“all unit tests in <code>crates/ragent-tui</code> pass”
Verify cmd	–	<code>cargo test -p ragent-tui</code>
Goal	“clean up the docs”	“every <code>.md</code> file in <code>docs/</code> has a one-line summary in <code>docs/INDEX.md</code> ”
Verify cmd	–	<code>python check_index.py</code>

Guidelines:

1. **State the end condition, not the journey.** The loop decides its own steps.
2. **Prefer a machine-checkable verify command.** Its exit status is what ends the loop; without one, the first tool-free answer ends it.
3. **Bound the blast radius.** Use `scope` to keep work inside a directory subtree and read-only to forbid writes to specific paths.
4. **Give a budget.** Unbounded loops burn tokens; set `max_steps` (dialog or config) so a stuck run ends with a diagnosable status.

6. The LoopSpec fields

LoopSpec (loop_state.rs) is the serialisable specification. `agent` and `goal` are required. `verify_cmd`, `scope`, `read_only`, `tool_set`, and `checkpoint_timeout_secs` are omitted from JSON when unset; `max_steps` and `cost_limit` serialise as `null` when unset (both may also be omitted on input); `checkpoints` is always present (`bool`) and must be supplied when deserialising a spec.

Field	Type	Default	Purpose
<code>agent</code>	<code>String</code>	<code>"general"</code>	Agent preset driving the loop
<code>goal</code>	<code>String</code>	<code>(empty)</code>	The success state; required to start
<code>verify_cmd</code>	<code>Option<String></code>	<code>None</code>	Verification command gating achievement; <code>None</code> means a no-tool-call response completes the loop directly
<code>scope</code>	<code>Vec<String></code>	<code>empty</code>	Scope-boundary globs; <code>empty</code> means the working directory with no extra restriction
<code>read_only</code>	<code>Vec<String></code>	<code>(empty)</code>	Read-only globs; writes to matching paths are denied and become observations
<code>tool_set</code>	<code>Vec<String></code>	<code>(empty)</code>	Restricted tool surface; <code>empty</code> means all tools
<code>max_steps</code>	<code>Option<u32></code>	<code>None</code>	Step budget; <code>None</code> resolves to the agent preset's <code>max_steps</code> , else the config default (512)
<code>cost_limit</code>	<code>Option<u64></code>	<code>None</code>	Token-cost gate (accumulated input + output); <code>None</code> resolves to the config value, which defaults to no gate
<code>checkpoints</code>	<code>bool</code>	<code>true</code>	Force destructive-action checkpoints even when an allow rule would auto-approve

Field	Type	Default	Purpose
checkpoint_timeout_secs	Option<u32>	None	Per-run checkpoint-prompt timeout override (seconds); None resolves to the loop.checkpoint_timeout_secs config default (120). Set by the one-shot --timeout flag.

Enforcement predicates on the spec:

- `allows_tool(tool)` – true when no tool set is configured, else membership.
- `path_in_scope(path)` – true when no scope is configured, else glob match.
- `path_is_read_only(path)` – glob match against `read_only`.
- `deny_reason(tool, input) -> Option<String>` – the full restriction check (Section 9).

Invalid glob patterns are silently skipped at compile time (they can match nothing, which fails closed).

7. Stop conditions and termination statuses

A loop run ends when exactly one stop condition fires; the first one wins and the tracker becomes stopped, after which no further stage may run.

Condition	Trigger	Published status
Goal achieved	Model responds without tool calls (plain goal loop)	completed
Verification passed	Verify command exits 0	completed (with the verification summary attached)
Step budget reached	<code>max_steps</code> iterations consumed	budget_exhausted
Token budget reached	accumulated input + output tokens $\geq$ <code>cost_limit</code>	budget_exhausted
Verification failing with budget breached	verify keeps failing once the budget is spent	budget_exhausted
Unrecoverable error	Provider/client failure during turn setup or the LLM stream	error
Consecutive recoverable failures exceeded	more than 3 (retry allowance) consecutive recoverable failures	error
User interrupt	Esc during a loop run	interrupted

Note the naming: the success variant is called `GoalAchieved` internally, but the published status string is `completed` – there is no `goal_achieved` status.

The TUI renders a status-aware banner from `loop_termination_banner(status, iterations)`:

Status	Banner
completed	Goal loop completed - {N} iteration(s)
error	Goal loop failed - {N} iteration(s)
budget_exhausted	Goal loop budget exhausted - {N} iteration(s)
interrupted	Goal loop interrupted - {N} iteration(s)
other	• Goal loop {other} - {N} iteration(s)

When present, the message window also appends `Verification: {outcome}` and `Reason: {why}`; the log panel records one line `loop terminated · {status} · {N} iteration(s)`.

8. The verification gate

The verification gate runs only when the active `LoopSpec` carries a `verify_cmd`. When the model produces a response without tool calls, the loop runs the command instead of accepting the answer:

- **Pass (exit 0)** – the loop ends with status `completed` and the captured verification output summary is attached to `Event::LoopTerminated`.
- **Fail, steps remaining** – the failure output is formatted into an observation (`verification_failure_observation`) and appended as the next observation in the chat history; the loop continues so the model can act on the failure output. A success is also recorded at this point, which *resets the consecutive-failure counter* so unrelated recoverable failures do not burn the retry allowance.
- **Fail, budget breached** – the run ends as `budget_exhausted`; there is no opportunity to retry the verification.

Execution details (mirroring the bash tool’s shell handling):

Property	Value
Execution	Shell command run inside the loop’s working directory
Timeout	600 seconds (<code>VERIFICATION_TIMEOUT</code>); a timeout counts as a failure
Output capture	head 8000 chars + tail 2000 chars, ASCII-truncated
Spawn error	Counts as failure (e.g. the program does not exist)
Status description	<code>exit code {N}</code> style reasons, mirroring the bash tool

9. Restrictions: tool set, read-only, scope

Before every tool execution during a loop run, `deny_reason` evaluates the restriction layers in a fixed order. A `None` result means the call proceeds to the normal permission layer (allow/deny/ask rules, YOLO, etc.).

1. **Tool-set restriction (FR-008/009)** – cheapest, most structural check. Mandatory safety tools are always exempted. Message: scope violation: tool out of scope – tool '<t>' is outside the loop's tool set (allowed: ...).
2. **Read-only constraint (FR-021)** – for bash, per-sub-command write heuristics plus redirection targets are inspected (e.g. sed -i, any > redirection); for other tools, LOOP_WRITE_TOOLS membership crossed with path globs. Message: constraint violation: path '<p>' is read-only for this loop (read-only constraints: ...).
3. **Scope boundaries (FR-022)** – every extracted path must be in scope. Message: scope violation: path '<p>' is outside the loop's scope boundaries (...).

A denial becomes a tool-error observation; the loop continues, and the model sees why the call was rejected.

Supporting constants:

- LOOP_WRITE_TOOLS – mutating tools, including legacy aliases: write, create, edit, multi_edit, multiedit, patch, apply_patch, append_to_file, update_file, write_file, rm, move_file, copy_file, make_directory, memory_store, memory_forget, memory_replace, memory_write, office_write, libre_write, pdf_write.
- LOOP_ALWAYS_ALLOWED_TOOLS – mandatory safety tools that stay available even under a restricted tool set: think, ask_user, agent_complete, model_info, memory_store, memory_recall, memory_forget, conversation_search, session_search.
- BASH_WRITE_COMMANDS – mutating shell commands (rm, mv, cp, mkdir, touch, truncate, shred, dd, tee, chmod, chown, chattr, install, ...); sed is handled specially for its in-place -i flag; git and package managers are deliberately excluded and left to the checkpoint layer.

10. Checkpoints and rollback

- **Forced checkpoints.** With checkpoints: true (the default), a *destructive* tool call inside a loop forces an interactive checkpoint prompt even when an allow rule (or YOLO mode) would have auto-approved it. The checkpoint waits checkpoint_timeout_secs (default 120); a timeout is treated as denial – the intervention defaults to the safe choice.
- **Pre-loop snapshot.** start_loop records the git state for the workspace immediately. The file snapshot itself is armed lazily at the first-write safe point (ensure_pre_loop_capture), so read-only loops never pay for one. If the workspace is not inside a git repository, a notice is published: loop workspace is not inside a git repository (FR-018); rollback will be snapshot-only.
- **Change summary and rollback offer.** When a loop that made workspace changes terminates, the processor publishes a change summary and the TUI offers:

```
Loop changes (completed, 3 iterations): 2 modified, 0 created, 0 deleted - +12 -3 lines added
Roll back to the pre-loop snapshot? Enter: roll back, Esc: keep changes
```

The status bar shows loop ended – Enter: roll back to pre-loop snapshot, Esc: keep changes. A no-write loop never publishes a change summary and never offers a rollback.

Rollback decision keys

While the offer is pending, the key interceptor consumes **every** key:

- Enter – confirm: restore the pre-loop snapshot. Log: rollback accepted · restoring {N} file(s).
- Esc – decline: keep all loop changes; pending captures are cleared. Message: Rollback declined – all changes made by the loop are kept.

If the restore reports nothing to restore, the status shows no pending capture (nothing to restore); on error the capture is kept for retry.

11. Configuration options

The loop section of `ragent.json` configures the goal-loop feature. It is **exclusively** the `/loop` feature's section – it does not tune the core per-turn loop (see `docs/howtos/reactagent.md` for `agent_perf.* /stream.*`).

```
{
  "loop": {
    "max_steps": 512,
    "cost_limit": 200000,
    "error_retry_allowance": 3,
    "checkpoints": true,
    "checkpoint_timeout_secs": 120
  }
}
```

Key	Type	Default	Effect
<code>loop.max_steps</code>	<code>u32</code>	512	Maximum loop iterations; exceeding it stops the loop with status <code>budget_exhausted</code>
<code>loop.cost_limit</code>	<code>Option<u64></code>	None	Maximum accumulated tokens (input + output) before the loop stops with status <code>budget_exhausted</code> ; None = no gate
<code>loop.error_retry_allowance</code>	<code>u8</code>	3	Consecutive recoverable-failure tolerance; the loop stops with status <code>error</code> once consecutive failures exceed this count
<code>loop.checkpoints</code>	<code>bool</code>	<code>true</code>	Force destructive-action checkpoints even when allow rules would auto-approve

Key	Type	Default	Effect
<code>loop.checkpoint_timeout</code>	<code>secs</code>	120	Seconds a forced checkpoint prompt waits before timeout counts as denial

Precedence rules

1. An explicit `max_steps` on the loop spec (dialog) wins over everything.
2. Otherwise the per-agent `agent.<name>.max_steps` override applies, if set.
3. Otherwise the `loop.max_steps` config default (512) applies.

The same precedence applies to the cost limit (spec value > config value, with `None` meaning “no gate” when nothing is configured).

Merge semantics

Config layers are deep-merged; overlay fields that differ from the compiled defaults win. `checkpoints` is monotonic during merge: once it is `false`, it stays `false`.

The serialized config omits the loop section entirely when every field is at its default value.

12. Commands and keyboard surface

Slash commands

Command	Effect
<code>/loop</code>	Open the interactive loop setup dialog (restores a cancelled draft if any)
<code>/loop <agent> <goal text...></code>	Start the loop immediately with the documented one-shot defaults
<code>/cancel</code>	Cancel in-flight processing
<code>/undo</code>	Snapshot rollback outside the loop flow
<code>/cost</code>	Run-cost summary for the session
<code>/autopilot on [--max-tokens N]</code> <code>[--max-time N]</code>	Autonomous continuation with its own budgets (separate feature)

Keyboard

Key	Effect
<code>Esc</code> (during a loop run)	Raise the loop interrupt – the loop stops after the current stage, status <code>interrupted</code>
<code>Enter</code> (while a rollback offer is pending)	Restore the pre-loop snapshot
<code>Esc</code> (while a rollback offer is pending)	Keep the loop’s changes

Programmatic surface (library)

Function	Behaviour
<code>SessionProcessor::start_loop(&self, session_id, spec)</code>	Register the spec + tracker, record git state, arm interrupt flag
<code>request_loop_interrupt(session_id)</code>	Set the interrupt flag; returns whether a loop was active
<code>rollback_loop(session_id)</code>	Restore the pre-loop capture; <code>Ok(false)</code> when nothing is pending
<code>clear_loop(session_id)</code>	Remove tracker, spec, interrupt flag, and capture
<code>clear_loop_captures()</code>	Clear all pending pre-loop captures (rollback decline)

There is **no HTTP endpoint** for starting a loop; `/loop` is driven from the TUI. (`POST /sessions/{id}/messages` drives plain turns, and autopilot is the HTTP/TUI autonomous-continuation feature.)

13. Worked examples

13.1 Drive the test suite to green

Interactive dialog:

1. `/loop`
2. `agent: coder`
3. `goal: cargo test -p ragent-tui passes with zero failures`
4. `verify cmd: cargo test -p ragent-tui`
5. `max steps: 40`
6. Enter

The loop edits code, runs the tool, receives failures as observations, and iterates until the verify command exits 0 or the budget is spent.

One-shot equivalent (no verify command, so the first tool-free answer ends the loop):

```
/loop coder make cargo test -p ragent-tui pass with zero failures
```

13.2 Scoped refactor with protected paths

Dialog values:

- `agent: coder`
- `goal: rename crate ragent-server to ragent-http across the workspace`
- `scope: crates/**, Cargo.toml`
- `read-only: specs/**, docs/**`
- `max steps: 60`
- `checkpoints: on`

Writes inside `specs/` or `docs/` are denied with a `constraint violation` observation; the loop can read them but not modify them.

13.3 Read-only research sweep

Dialog values:

- agent: general
- goal: produce a migration risk list for every TODO in src/
- read-only: ** (nothing writable)
- max steps: 30

With nothing writable the loop cannot mutate the workspace; the final tool-free answer is the deliverable and ends the loop with completed.

13.4 Budgeted doc-coverage pass (curl-driven session, then TUI loop)

Start a session and drive a plain turn over HTTP (the loop itself is started in the TUI):

```
curl -s -X POST http://127.0.0.1:3000/sessions \  
  -H "Authorization: Bearer $RAGENT_TOKEN" \  
  -H "Content-Type: application/json" \  
  -d '{"directory": "."}'  
  
curl -s -X POST http://127.0.0.1:3000/sessions/<id>/messages \  
  -H "Authorization: Bearer $RAGENT_TOKEN" \  
  -H "Content-Type: application/json" \  
  -d '{"content": "list every public function missing doc comments in crates/ragent-agent/src/ses"
```

Then, in the TUI against the same session, program the fix loop:

```
/loop coder add doc comments to every function in the missing-docs list
```

13.5 CLI one-shot, then inspect the outcome

```
ragent run --agent coder --maxsteps 50 "audit specs/agentloop against its FRs and report gaps"
```

- --maxsteps (single word) caps the agentic loop steps.
- --yes (alias --no-prompt) auto-approves permission prompts – note that forced loop checkpoints still interrupt destructive actions when checkpoints are on.
- ragent serve --addr 127.0.0.1:9100 starts the HTTP API (there is no --port flag).

14. Events and telemetry

Two workspace events carry the loop lifecycle to every subscriber (TUI, SSE):

Event::LoopTerminated – published exactly once when a stop condition fires. Fields:

Field	Meaning
session_id	Session that ran the loop
status	Termination label: completed, error, budget_exhausted, interrupted
iterations	Iterations consumed

Field	Meaning
verification	Verification-command output summary when one ran
reason	Human-readable explanation for non-success terminations

Event::LoopChangeSummary – published when a loop that changed the workspace terminates. Fields:

Field	Meaning
status	Mirrors <code>LoopTerminated.status</code>
iterations	Iterations consumed
files_modified, files_created, files_deleted	Change counts
diffstat	e.g. "+12 -3 lines across 4 files"
files	Sorted paths, relative to the workspace root

Telemetry is recorded exactly once per loop run (`loop_telemetry_recorded` flag): the loop duration and iteration count plus the per-run tool-call tally are written through the telemetry recorder when the run terminates with a known start time. Starting a new loop resets the flag.

15. Related documentation

- `docs/howtos/reactagent.md` – the core per-turn ReACT agent loop (what runs underneath every `/loop` iteration), including `agent_perf.*`, `stream.*`, and compaction tuning.
- `TUI-QUICKSTART.md` – section “Goal-driven loops with `/loop`” for a quick-reference version of this guide.
- `specs/agentloop/SPEC.md` – the requirements specification (FR-001..FR-025) this feature implements.