

ragent Repository Reverse-Engineering Manual

This guide explains how to use ragent's `/spec reverse` command to reverse-engineer the purpose, architecture, and function of public GitHub repositories. The command fetches a repository's metadata, root file tree, and README via the GitHub API, then asks the currently selected LLM model to generate a synthetic "creation prompt" — a detailed prompt that, if fed to a coding agent, would reproduce the repository from scratch.

The same functionality is available as the `ragent reverse` CLI subcommand.

Scope: `/spec reverse` slash commands, the `/new scaffold` flags (`--language` / `--type` / `--stack`), the scaffold-only `--folder` and `--github` / `--gitlab` flags, the `--create` and `--depth` flags, GitHub and GitLab API interaction, synthetic prompt generation, project scaffolding, and spec chaining. For spec management commands, see `docs/howtos/spec.md`. For research commands, see `docs/howtos/research.md`.

1. Purpose and capabilities

The `/spec reverse` command answers a common question when encountering an unfamiliar repository: **"What does this project do, and how would I build something like it?"**

Instead of manually browsing a repo's file tree, reading the README, and guessing at the architecture, `/spec reverse` automates the entire process:

- **Fetches repository metadata** via the GitHub REST API — description, primary language, license, star count, topics, creation date, and homepage URL.
- **Fetches the root file tree** — the top-level directory listing of the repository, including file and directory names.
- **Fetches the README** — the full rendered README content (when available).
- **Generates a synthetic creation prompt** — passes all gathered information to the currently selected LLM model and asks it to produce a comprehensive prompt that describes how to recreate the repository from scratch, including architecture, key modules, technology choices, and implementation order.
- **Optionally constrains the target project shape** — the `/new scaffold` flags (`--language`, `--type`, and the optional `--stack`) let you specify a target language, app type, and framework stack so the generated prompt is tailored to that shape rather than the repository's original languages.
- **Optionally scaffolds the project** — with the scaffold flags present, `--folder <path>` runs the same `/new engine /spec govcreeate` uses to scaffold a real project (default: the current directory), and `--github` / `--gitlab` create a private remote and push the initial commit.
- **Optionally chains into spec creation** — the `--create` flag feeds the generated prompt directly into `/spec create`, auto-generating a formal specification from the reverse-engineered prompt.

What it does

- Fetches public repository data using the GitHub API (a GitHub token is required — see Section 8).
- Derives a comprehensive creation prompt that captures the repository's architecture, module structure, key features, and technology decisions.
- Streams the generated prompt into the chat window for review.
- Can chain into the spec management system to produce a formal SPEC.md.

What it does not do

- It does not clone the repository — it uses the GitHub API only.
 - It does not read source files beyond the root tree listing — it does not fetch or parse individual source files.
 - It does not work with private repositories unless GITHUB_TOKEN is set with appropriate access.
 - It does not generate code — it generates a prompt that *describes* how to build code. (With scaffold flags it emits the same runnable hello-world starter `/new` produces, not a re-implementation of the source repo.)
-

2. Quick start

Open ragent and run:

```
/spec reverse thrivethrough/omitme
```

The TUI will:

1. Fetch the repository metadata, root file tree, and README from the GitHub API.
2. Pass all gathered information to the currently selected LLM model.
3. Ask the model to generate a synthetic creation prompt.
4. Display the generated prompt in the chat window.

To chain directly into spec creation:

```
/spec reverse thrivethrough/omitme --create omitme-clone
```

This fetches the repository, generates the creation prompt, and immediately passes it to `/spec create` to produce a formal specification under `specs/omitme-clone/`.

3. Command syntax

3.1 `/spec reverse <owner/repo>`

Fetch a public GitHub repository and generate a synthetic creation prompt.

```
/spec reverse <repo> [--language <lang> --type <type> [--stack <name>]] [--create <name>] [--dept  
/spec reverse <repo> --language <lang> --type <type> [--stack <name>] [--folder <path>] [--github
```

Arguments

Argument	Required	Description
<repo>	Yes	Repository identifier. Accepts owner/repo, github:owner/repo, github:<github-url>, gitlab:namespace/project, gitlab:host/namespace/project, HTTPS URLs (https://github.com/owner/repo, https://gitlab.com/ns/proj), and SSH URLs (git@github.com:owner/repo.git)

Flags

Flag	Required	Description
--language <lang>	No	Target language for the generated prompt — the same values /new accepts (rust, python, go, typescript, ...). Must be supplied with --type
--type <type>	No	Target app type — the same values /new accepts: library, cmdline, tui, gui, webapp. Must be supplied with --language
--stack <name>	No	Optional framework stack layer — the same values /new accepts (axum, warp, raylib, gtk4, ratatui). Requires --language and --type
--create <name>	No	Chain into /spec create using the generated prompt. With --folder, the spec is written to <folder>/specs/<name>/
--depth <N>	No	Directory levels to fetch from the tree (1–10, default 1). See Section 10.
--folder <path>	No	With scaffold flags, scaffold the project in <path> (created if missing; default: the current directory). See Section 6.1
--github	No	With scaffold flags, create a private GitHub repository, set it as origin, and push the initial commit. Mutually exclusive with --gitlab
--gitlab	No	With scaffold flags, create a private GitLab repository, set it as origin, and push the initial commit. Mutually exclusive with --github

--language, --type, and --stack are the /new scaffold flags, parsed and validated by the same /new parser, so they take exactly the same values and have the same purpose: they steer the generated prompt

towards a target project shape. `--language` and `--type` are all-or-nothing — supplying one without the other (or a lone `--stack`) is a usage error, exactly as with `/new`. Run `/new help` for the full accepted-value lists.

`--folder`, `--github`, and `--gitlab` only take effect when the scaffold flags are present — supplying any of them without `--language` and `--type` is a usage error.

`<repo>` must be the first argument after `reverse`. An invocation that starts with a flag (e.g. `/spec reverse --language rust --type cmdline`) is a usage error that names the offending token, not a silent no-op.

3.2 Input formats

The command accepts several input formats:

```
# owner/repo shorthand
/spec reverse thrivethrough/omitme
```

```
# Full GitHub URL
/spec reverse https://github.com/thrivethrough/omitme
```

```
# With a target scaffold (language + app type)
/spec reverse thrivethrough/omitme --language rust --type cmdline
```

```
# With spec creation
/spec reverse thrivethrough/omitme --create omitme-clone
```

```
# With both
/spec reverse thrivethrough/omitme --language rust --type cmdline --create omitme-rust-port
```

```
# Scaffold a project from the reverse-engineered prompt
/spec reverse thrivethrough/omitme --language rust --type cmdline --folder ./omitme-clone
```

```
# Scaffold it and create a private GitHub repository
/spec reverse thrivethrough/omitme --language rust --type cmdline --folder ./omitme-clone --github
```

3.3 `/spec reverse help`

Show the command reference.

```
/spec reverse help
```

4. What the command gathers

When you run `/spec reverse <owner/repo>`, the system fetches three pieces of information from the GitHub REST API:

4.1 Repository metadata

Fetches from GET `/repos/{owner}/{repo}`:

- **Name and full name** — e.g. `thrivethrough/omitme`
- **Description** — the repository’s short description
- **Primary language** — the dominant programming language
- **License** — the detected license (e.g. MIT, Apache-2.0)
- **Star count** — number of GitHub stars
- **Topics** — repository topics/tags
- **Creation date** — when the repository was created
- **Homepage URL** — the repository’s homepage (if set)
- **Default branch** — e.g. `main` or `master`

4.2 Root file tree

Fetches from `GET /repos/{owner}/{repo}/contents:`

The top-level directory listing of the repository, including all files and directories at the root level. This gives the LLM a structural overview of the project layout — enough to infer the architecture, build system, and module organisation without reading individual files.

4.3 README content

Fetches from `GET /repos/{owner}/{repo}/readme:`

The full README content, rendered as markdown. The README typically contains the project overview, installation instructions, usage examples, and architecture notes — the richest single source of information about a repository’s purpose and design.

5. The synthetic creation prompt

The core output of `/spec reverse` is a **synthetic creation prompt** — a natural-language description of how to recreate the repository from scratch.

The LLM receives the repository metadata, root file tree, and README, and is instructed to produce a prompt that:

1. **Summarises the project’s purpose** — what the project does and who it is for.
2. **Describes the architecture** — the high-level module structure, key components, and how they interact.
3. **Identifies the technology stack** — languages, frameworks, libraries, build tools, and runtime dependencies.
4. **Outlines the implementation order** — a suggested sequence for building the project, starting with foundational modules and progressing to higher-level features.
5. **Highlights key design decisions** — notable architectural choices, patterns, and trade-offs visible from the file tree and README.

The generated prompt is designed to be fed to a coding agent (like `ragent` itself) to reproduce the repository’s functionality from scratch.

Example generated prompt structure

```
## Project: omitme - Privacy-first error tracking
```

Overview

omitme is a privacy-first error tracking service that captures application errors without collecting personally identifiable information (PII). The project is built in Rust and uses a client-server architecture with a WebSocket-based real-time event stream.

Architecture

The project consists of the following top-level modules:

- ``src/server/`` - Axum-based HTTP server handling event ingestion
- ``src/client/`` - SDK libraries for Rust, TypeScript, and Python
- ``src/storage/`` - SQLite-backed event store with retention policies
- ``src/analytics/`` - Aggregation and dashboard query engine
- ``src/web/`` - Static dashboard frontend (HTML/CSS/JS)

Technology Stack

- **Language:** Rust (edition 2024)
- **Web framework:** Axum
- **Database:** SQLite via rusqlite
- **Real-time:** tokio-tungstenite (WebSockets)
- **Serialization:** serde + serde_json

Implementation Order

1. Core data model and SQLite schema (``src/storage/``)
2. Event ingestion API endpoint (``src/server/``)
3. WebSocket event stream (``src/server/``)
4. Client SDKs (``src/client/``)
5. Analytics and aggregation queries (``src/analytics/``)
6. Dashboard frontend (``src/web/``)

Key Design Decisions

- PII stripping happens client-side before transmission
- Events are stored with a configurable retention period
- The dashboard is a static site with no server-side rendering

6. The scaffold flags (`--language` / `--type` / `--stack`)

By default, the generated prompt reflects the repository's original technology stack. The `/new-style` scaffold flags — `--language <lang>`, `--type <type>`, and the optional `--stack <name>` — constrain the prompt to a specific target project shape, making the result useful for planning a port or rewrite. They are parsed and validated by the same `/new` parser, so they take exactly the same values `/new` accepts: `--language rust --type cmdline` or `--language rust --type tui --stack ratatui`. `--language` and `--type` must be supplied together; a lone `--stack` is rejected with the same usage error

/new produces. Run `/new help` for the full accepted-value lists.

How it works

When the scaffold flags are supplied, the LLM is instructed to:

- Describe how to recreate the repository's *functionality* using the specified language, app type, and stack.
- Map the original architecture to equivalent components in the target stack.
- Note where the target stack differs from the original and what adjustments are needed.

Examples

Port a Python project to a Rust command-line app:

```
/spec reverse example/flask-api --language rust --type cmdline
```

Port a Node.js project to Python:

```
/spec reverse example/express-server --language python --type cmdline
```

Port a Go project to a TypeScript web app:

```
/spec reverse example/go-microservice --language typescript --type webapp
```

Target a Rust TUI app on the ratatui stack:

```
/spec reverse example/legacy-java-app --language rust --type tui --stack ratatui
```

6.1 Scaffolding with `--folder`

With the scaffold flags present, `--folder <path>` does more than steer the prompt: it scaffolds a real project in `<path>` using the same engine `/new` and `/spec govcreate` use. The engine is shared, so the result is byte-for-byte what `/new` would produce for those flags — workspace artifacts (`.ragent/`, `specs/`, `log/`, `.gitignore`, `AGENTS.md`), the language layout and hello-world starter, `README.md` / `QUICKSTART.md` / `STATS.md` / `docs/`, a `git init` with an initial commit, and, with `--github` / `--gitlab`, a private remote set as `origin` plus a push.

`<path>` is created if it does not exist. When `--folder` is omitted, the current directory is scaffolded in place (mirroring `ragent new`). The target must be empty apart from `.ragent/`, `log/`, and `target/` — the same guard `/new` runs.

The scaffold step runs **before** prompt synthesis, so a refusal is reported immediately:

- a non-empty target folder reports `[err] target folder is not empty` with the blocking entries and the run continues to generate the prompt;
- a hosting failure (e.g. no GitHub token) reports the failing step and the local scaffold remains intact.

In either case nothing is scaffolded and the generated prompt is still produced, so the command is safe to re-run after fixing the cause.

When `--create <name>` is supplied alongside `--folder`, the chained spec is written into the scaffolded project (`<folder>/specs/<name>/`), not the invoking directory — see Section 7.

Examples:

```
# Scaffold a Rust command-line project, keeping the prompt steering
/spec reverse example/legacy-java-app --language rust --type cmdline --folder ./legacy-port

# Scaffold it and create a private GitLab repository
/spec reverse example/go-microservice --language python --type library --folder ./svc --gitlab
```

7. The `--create` flag

The `--create <name>` flag chains the reverse-engineering output into `/spec create`, auto-generating a formal specification from the reverse-engineered prompt.

How it works

1. `/spec reverse` fetches the repository and generates the synthetic creation prompt.
2. The generated prompt is passed as the feature description to `/spec create <name>` — with `--folder`, as `/spec create <name> --folder <path> <prompt>`, so the spec root is the scaffolded project's.
3. The spec system creates `specs/<name>/SPEC.md` (EARS requirements), `specs/<name>/PLAN.md` (implementation plan), and `specs/<name>/TESTPLAN.md` (manual test plan) under that root.

Examples

Reverse-engineer a repo and create a spec:

```
/spec reverse thrivethrough/omitme --create omitme-clone
```

Reverse-engineer with a tech constraint and create a spec:

```
/spec reverse thrivethrough/omitme --language rust --type cmdline --create omitme-rust-port
```

Reverse-engineer a large project into a spec:

```
/spec reverse tokio-rs/tokio --create tokio-study
```

What you get

After `--create` finishes, the spec directory contains:

```
specs/omitme-clone/
  SPEC.md          # EARS requirements derived from the creation prompt
  PLAN.md          # Implementation plan with task table
  TESTPLAN.md      # Manual test plan with test cases
```

With `--folder <path>`, the same tree is written to `<path>/specs/omitme-clone/` inside the scaffolded project.

You can then use the standard `/spec` commands to validate, plan, and implement the spec (see `docs/howtos/spec.md` for details).

8. API interaction

GitHub authentication

A GitHub token is a **prerequisite**: `/spec reverse` refuses to run without one and reports “No GitHub token configured. Run `/github login` to authenticate, then re-run `/spec reverse`.” The token is resolved via `/github login` (OAuth device flow, stored in `~/.ragent/github_token`) or the `GITHUB_TOKEN` environment variable. Unauthenticated GitHub API access allows only 60 requests per hour per IP; an authenticated token raises this to 5,000 requests per hour.

```
export GITHUB_TOKEN="ghp_your_token_here"
```

The token does not need any special scopes for reading public repositories. A fine-grained token with read access to public repositories is sufficient.

GitLab authentication

GitLab API calls require a personal access token. The token is resolved in priority order:

1. `GITLAB_TOKEN` environment variable
2. `ragent.json` configuration
3. Encrypted credential database (configured via `/gitlab setup`)

For self-hosted GitLab instances, the host is resolved in priority order: explicit host in the repo identifier (`gitlab:host/namespace/project`), then the `GITLAB_URL` environment variable, then `https://gitlab.com`.

```
export GITLAB_TOKEN="glpat-your_token_here"
export GITLAB_URL="https://gitlab.example.com"
```

Rate limiting

If the API rate limit is exceeded, the command will report an error. Wait for the rate limit window to reset (typically 1 hour) or set `GITHUB_TOKEN` for higher limits.

When using `--depth` for recursive tree fetch, each subdirectory level adds additional API calls (one per subdirectory). Use a token for deeper trees to avoid hitting the rate limit.

Error handling

Error	Cause	What to do
repository not found	Invalid owner/repo or private repo	Check the spelling; set <code>GITHUB_TOKEN</code> for private repos
rate limit exceeded	Too many API requests	Wait for reset or set <code>GITHUB_TOKEN</code>
README not found	Repository has no README	The command continues with metadata and tree only
network error	Connection issue	Check network connectivity and retry
gitlab token missing	No GitLab token configured	Set <code>GITLAB_TOKEN</code> env var or run <code>/gitlab setup</code>
invalid depth	<code>--depth</code> value out of range	Use a value between 1 and 10

Error	Cause	What to do
target folder is not empty	--folder target has entries other than .ragent/, log/, target/	Empty the folder or choose another --folder; the prompt is still generated
--folder requires --language and --type	--folder/--github/--gitlab used without scaffold flags	Add --language <lang> --type <type>
no GitHub/GitLab token for --github/--gitlab	Hosting requested without credentials	Run /github login or /gitlab setup; the local scaffold remains intact

Silent degradation on GitHub fetch failures: metadata/tree fetch failures on the GitHub path are swallowed (`unwrap_or_default()`), so an API error produces an empty context fed to the LLM rather than an error message — the repository not found / rate limit exceeded / network error messages only surface on the **GitLab** path. If the generated prompt looks thin, re-run after checking your token and rate limit.

9. CLI equivalents

The TUI `/spec reverse` command is the primary entry point; the equivalent `ragent reverse` CLI subcommand is not currently registered in the clap CLI (`run`, `serve`, `session`, `memory`, `auth`, `models`, `config`, `research` are the available subcommands). The examples below describe the intended CLI surface and will become accurate when the subcommand ships:

Basic reverse-engineering

```
ragent reverse thrivethrough/omitme
```

With technology constraint

```
ragent reverse thrivethrough/omitme --language rust --type cmdline
```

With spec creation

```
ragent reverse thrivethrough/omitme --create omitme-clone
```

Full GitHub URL

```
ragent reverse https://github.com/thrivethrough/omitme
```

Both flags

```
ragent reverse thrivethrough/omitme --language rust --type cmdline --create omitme-rust-port
```

Scaffold the project (and optionally a private hosting remote)

```
ragent reverse thrivethrough/omitme --language rust --type cmdline --folder ./omitme-clone --github
```

The CLI command prints the generated prompt to stdout, making it easy to pipe into other tools or save to a file:

Save the generated prompt to a file

```
ragent reverse thrivethrough/omitme > omitme-prompt.md
```

```
# Pipe through jq for structured output
ragent reverse thrivethrough/omitme --language rust --type cmdline 2>&1 | tee omitme-rust.md
```

10. The `--depth` flag

By default, `/spec reverse` fetches only the root-level file tree (depth 1). The `--depth <N>` flag controls how many levels of subdirectories are expanded.

Depth	What you get
1 (default)	Root-level files and directories only
2	Root + one level of subdirectories
3	Root + two levels of subdirectories
N (max 10)	Root + N-1 levels of subdirectories

Deeper trees give the LLM more context about the project's module structure, which improves the quality of the generated creation prompt — especially for large repositories with deep nesting. However, each additional level requires more API calls (one per subdirectory), so use a token for higher rate limits when fetching deep trees.

Examples

```
# Default depth (root only)
/spec reverse thrivethrough/omitme

# Two levels deep
/spec reverse thrivethrough/omitme --depth 2

# Maximum depth for a large project
/spec reverse torvalds/linux --depth 5 --language rust --type cmdline
```

11. End-to-end examples

Example 1: Understand an unfamiliar project

You find an interesting repository and want to understand what it does:

```
/spec reverse thrivethrough/omitme
```

The generated prompt gives you a comprehensive overview of the project's purpose, architecture, and technology stack — without needing to browse the code yourself.

Example 2: Plan a port to a different language

You want to port a Python project to Rust:

```
/spec reverse example/flask-api --language rust --type cmdline
```

The generated prompt describes how to recreate the Flask API's functionality using Rust, mapping each component to its Rust equivalent (e.g. Flask → Axum, SQLAlchemy → SQLx, Celery → Tokio tasks).

Example 3: Create a spec from a reference implementation

You want to create a formal spec based on an existing open-source project:

```
/spec reverse tokio-rs/tokio --create tokio-study
```

This generates a spec under `specs/tokio-study/` with EARS requirements, an implementation plan, and a test plan — all derived from the reverse-engineered creation prompt.

Example 4: Study a well-architected project

You want to study how a well-known project is structured:

```
/spec reverse burntsushi/regex
```

The generated prompt breaks down the regex crate's architecture, module structure, and key design decisions, giving you a learning roadmap.

Example 5: Plan a rewrite with a specific stack

You want to rewrite a legacy application using a modern stack:

```
/spec reverse example/legacy-monolith --language rust --type cmdline --stack axum
```

The quotes are required here: an unquoted value would stop at the first space. The generated prompt maps the legacy monolith's functionality to a modern Rust microservice architecture with specific crate recommendations.

Example 6: Reverse-engineer and immediately implement

Combine `/spec reverse` with `/spec impl` for a full pipeline:

```
# Step 1: Reverse-engineer and create a spec
```

```
/spec reverse thrivethrough/omitme --language rust --type cmdline --create omitme-rust
```

```
# Step 2: Validate the generated spec
```

```
/spec validate omitme-rust
```

```
# Step 3: Preview the implementation plan
```

```
/spec impl omitme-rust --dry-run
```

```
# Step 4: Implement
```

```
/spec impl omitme-rust
```

Example 7: Compare two implementations

Reverse-engineer two competing projects to compare their approaches:

```
/spec reverse project-a/repo --create study-a
```

```
/spec reverse project-b/repo --create study-b
```

Then compare the generated specs to understand the architectural differences between the two projects.

Example 8: Generate a prompt for a specific use case

You want a creation prompt focused on the testing infrastructure of a project:

```
/spec reverse example/well-tested-app --language rust --type cmdline
```

The generated prompt emphasises how to recreate the project’s testing approach using the specified testing tools.

Example 9: Reverse-engineer a GitLab repository

```
# Self-hosted GitLab with nested namespace
```

```
/spec reverse gitlab:gitlab.example.com/group/subgroup/project --depth 3
```

```
# GitLab.com with spec creation
```

```
/spec reverse gitlab:my-namespace/my-project --create my-project-clone
```

```
# Full GitLab URL
```

```
/spec reverse https://gitlab.com/my-namespace/my-project --language python --type cmdline --creat
```

12. Tips for good results

- **Use the scaffold flags for targeting.** `--language rust --type cmdline` is good; add `--stack ratatui` (or `--stack axum`) when you want the prompt to reach for a specific framework. Values come from the `/new` registry, so an unknown value is reported as a usage error before any fetch happens.
 - **Review the generated prompt before acting.** The synthetic prompt is a starting point — review it for accuracy and adjust before feeding it to a coding agent.
 - **Use `--create` for structured workflows.** Chaining into `/spec create` gives you a formal spec, plan, and test plan — much more actionable than a raw prompt.
 - **Use `--folder` to get a runnable starting point.** Scaffolding emits the same starter `/new` produces (not a re-implementation of the source repo), which is often a better base than hand-writing the first files.
 - **Set `GITHUB_TOKEN` for frequent use.** A token is required to run the command at all; the unauthenticated API allows only 60 requests/hour, authenticated is 5,000/hour.
 - **Use full URLs for clarity.** `/spec reverse https://github.com/owner/repo` is unambiguous; `/spec reverse owner/repo` is faster to type.
 - **Combine with `/spec` commands.** After `--create`, use `/spec validate`, `/spec plan`, `/spec tasks`, and `/spec impl` for a complete spec-driven workflow.
 - **Use the scaffold flags for port planning.** When planning a port, always specify the target language and app type so the prompt maps to the right technologies.
 - **Try different models.** The generated prompt quality depends on the selected LLM model. Try different models (e.g. Claude, GPT-4) to compare output quality.
 - **Study well-architected projects.** Use `/spec reverse` on projects you admire to learn about their architecture and design patterns.
-

13. Troubleshooting

Symptom	Likely cause	What to do
repository not found	Invalid owner/repo or private repo	Check spelling; set GITHUB_TOKEN for private repos
rate limit exceeded	Too many unauthenticated API calls	Set GITHUB_TOKEN environment variable
README not found	Repository has no README	The command continues with metadata and tree only
Empty or short prompt	LLM model not configured or down	Check provider setup with /models
network error	Connection issue	Check network connectivity and retry
invalid repository format	Malformed input	Use owner/repo or a full GitHub URL
spec already exists	--create target name already used	Choose a different name or delete the existing spec
Prompt mentions wrong technologies	--language/--type/--stack not specified or too vague	Specify a more precise scaffold (e.g. --language rust --type cmdline --stack axum)

14. Workflow integration

The /spec reverse command integrates with ragent's spec management system to provide a complete reverse-engineering-to-implementation pipeline:

```
/spec reverse <owner/repo>
```

```
scaffold flags + --folder/--github/--gitlab?
```

```
/new scaffold engine
(emit + git init +
optional remote)
```

```
refusal reported, run continues
```

```
GitHub API fetch
(metadata + tree
+ README)
```

```
LLM generates
creation prompt
```

```
                --create <name>
Display prompt      /spec create
in chat window      generates SPEC.md,
                    PLAN.md, TESTPLAN.md
```

```
                /spec validate
                /spec plan
                /spec impl
```

Full pipeline example

```
# 1. Reverse-engineer a project
/spec reverse thrivethrough/omitme --language rust --type cmdline --create omitme-rust

# 2. Validate the generated spec
/spec validate omitme-rust

# 3. Generate task list
/spec tasks omitme-rust

# 4. Preview implementation
/spec impl omitme-rust --dry-run

# 5. Implement
/spec impl omitme-rust
```

This pipeline takes you from an unfamiliar GitHub repository to a running implementation in five commands.

End of manual.