

ragent Repository Reverse-Engineering Manual

This guide explains how to use ragent's `/reverse` command to reverse-engineer the purpose, architecture, and function of public GitHub repositories. The command fetches a repository's metadata, root file tree, and README via the GitHub API, then asks the currently selected LLM model to generate a synthetic "creation prompt" — a detailed prompt that, if fed to a coding agent, would reproduce the repository from scratch.

The same functionality is available as the `ragent reverse` CLI subcommand.

Scope: `/reverse` slash commands, the `--tech` and `--create` flags, GitHub API interaction, synthetic prompt generation, and spec chaining. For spec management commands, see `docs/howtos/spec.md`. For research commands, see `docs/howtos/research.md`.

1. Purpose and capabilities

The `/reverse` command answers a common question when encountering an unfamiliar repository: **"What does this project do, and how would I build something like it?"**

Instead of manually browsing a repo's file tree, reading the README, and guessing at the architecture, `/reverse` automates the entire process:

- **Fetches repository metadata** via the GitHub REST API — description, primary language, license, star count, topics, creation date, and homepage URL.
- **Fetches the root file tree** — the top-level directory listing of the repository, including file and directory names.
- **Fetches the README** — the full rendered README content (when available).
- **Generates a synthetic creation prompt** — passes all gathered information to the currently selected LLM model and asks it to produce a comprehensive prompt that describes how to recreate the repository from scratch, including architecture, key modules, technology choices, and implementation order.
- **Optionally constrains the technology stack** — the `--tech` flag lets you specify a target technology stack so the generated prompt is tailored to that stack rather than the repository's original languages.
- **Optionally chains into spec creation** — the `--create` flag feeds the generated prompt directly into `/spec create`, auto-generating a formal specification from the reverse-engineered prompt.

What it does

- Fetches public repository data using the GitHub API (a GitHub token is required — see Section 8).
- Derives a comprehensive creation prompt that captures the repository's architecture, module structure, key features, and technology decisions.
- Streams the generated prompt into the chat window for review.
- Can chain into the spec management system to produce a formal SPEC.md.

What it does not do

- It does not clone the repository — it uses the GitHub API only.
- It does not read source files beyond the root tree listing — it does not fetch or parse individual source files.
- It does not work with private repositories unless `GITHUB_TOKEN` is set with appropriate access.

- It does not generate code — it generates a prompt that *describes* how to build code.

2. Quick start

Open ragent and run:

```
/reverse thrivesthrough/omitme
```

The TUI will:

1. Fetch the repository metadata, root file tree, and README from the GitHub API.
2. Pass all gathered information to the currently selected LLM model.
3. Ask the model to generate a synthetic creation prompt.
4. Display the generated prompt in the chat window.

To chain directly into spec creation:

```
/reverse thrivesthrough/omitme --create omitme-clone
```

This fetches the repository, generates the creation prompt, and immediately passes it to `/spec create` to produce a formal specification under `specs/omitme-clone/`.

3. Command syntax

3.1 `/reverse <owner/repo>`

Fetch a public GitHub repository and generate a synthetic creation prompt.

```
/reverse <repo> [--tech <stack>] [--create <name>] [--depth <N>]
```

Arguments

Argument	Required	Description
<code><repo></code>	Yes	Repository identifier. Accepts <code>owner/repo</code> , <code>github:owner/repo</code> , <code>github:<github-url></code> , <code>gitlab:namespace/project</code> , <code>gitlab:host/namespace/project</code> , HTTPS URLs (<code>https://github.com/owner/repo</code> , <code>https://gitlab.com/ns/proj</code>), and SSH URLs (<code>git@github.com:owner/repo.git</code>)

Flags

Flag	Required	Description
<code>--tech <stack></code>	No	Constrain the generated prompt to a specific technology stack (e.g. rust, python, typescript). Values containing spaces must be quoted — <code>--tech "Next.js + Rails"</code> is captured whole; unquoted values are single tokens, so hyphenated forms like <code>rust-axum-sqlx</code> also work
<code>--create <name></code>	No	Chain into <code>/spec create</code> using the generated prompt, creating a spec under <code>specs/<name>/</code>
<code>--depth <N></code>	No	Directory levels to fetch from the tree (1–10, default 1). See Section 10.

3.2 Input formats

The command accepts several input formats:

```
# owner/repo shorthand
/reverse thrivesthrough/omitme

# Full GitHub URL
/reverse https://github.com/thrivethrough/omitme

# With technology constraint
/reverse thrivesthrough/omitme --tech rust

# With spec creation
/reverse thrivesthrough/omitme --create omitme-clone

# With both flags
/reverse thrivesthrough/omitme --tech rust --create omitme-rust-port
```

3.3 `/reverse help`

Show the command reference.

```
/reverse help
```

4. What the command gathers

When you run `/reverse <owner/repo>`, the system fetches three pieces of information from the GitHub REST API:

4.1 Repository metadata

Fetches from `GET /repos/{owner}/{repo}`:

- **Name and full name** — e.g. `thrivethrough/omitme`
- **Description** — the repository’s short description
- **Primary language** — the dominant programming language
- **License** — the detected license (e.g. MIT, Apache-2.0)
- **Star count** — number of GitHub stars
- **Topics** — repository topics/tags
- **Creation date** — when the repository was created
- **Homepage URL** — the repository’s homepage (if set)
- **Default branch** — e.g. `main` or `master`

4.2 Root file tree

Fetches from `GET /repos/{owner}/{repo}/contents:`

The top-level directory listing of the repository, including all files and directories at the root level. This gives the LLM a structural overview of the project layout — enough to infer the architecture, build system, and module organisation without reading individual files.

4.3 README content

Fetches from `GET /repos/{owner}/{repo}/readme:`

The full README content, rendered as markdown. The README typically contains the project overview, installation instructions, usage examples, and architecture notes — the richest single source of information about a repository’s purpose and design.

5. The synthetic creation prompt

The core output of `/reverse` is a **synthetic creation prompt** — a natural-language description of how to recreate the repository from scratch.

The LLM receives the repository metadata, root file tree, and README, and is instructed to produce a prompt that:

1. **Summarises the project’s purpose** — what the project does and who it is for.
2. **Describes the architecture** — the high-level module structure, key components, and how they interact.
3. **Identifies the technology stack** — languages, frameworks, libraries, build tools, and runtime dependencies.
4. **Outlines the implementation order** — a suggested sequence for building the project, starting with foundational modules and progressing to higher-level features.
5. **Highlights key design decisions** — notable architectural choices, patterns, and trade-offs visible from the file tree and README.

The generated prompt is designed to be fed to a coding agent (like `ragent` itself) to reproduce the repository’s functionality from scratch.

Example generated prompt structure

```
## Project: omitme - Privacy-first error tracking
```

Overview

omitme is a privacy-first error tracking service that captures application errors without collecting personally identifiable information (PII). The project is built in Rust and uses a client-server architecture with a WebSocket-based real-time event stream.

Architecture

The project consists of the following top-level modules:

- ``src/server/`` - Axum-based HTTP server handling event ingestion
- ``src/client/`` - SDK libraries for Rust, TypeScript, and Python
- ``src/storage/`` - SQLite-backed event store with retention policies
- ``src/analytics/`` - Aggregation and dashboard query engine
- ``src/web/`` - Static dashboard frontend (HTML/CSS/JS)

Technology Stack

- **Language:** Rust (edition 2024)
- **Web framework:** Axum
- **Database:** SQLite via rusqlite
- **Real-time:** tokio-tungstenite (WebSockets)
- **Serialization:** serde + serde_json

Implementation Order

1. Core data model and SQLite schema (``src/storage/``)
2. Event ingestion API endpoint (``src/server/``)
3. WebSocket event stream (``src/server/``)
4. Client SDKs (``src/client/``)
5. Analytics and aggregation queries (``src/analytics/``)
6. Dashboard frontend (``src/web/``)

Key Design Decisions

- PII stripping happens client-side before transmission
- Events are stored with a configurable retention period
- The dashboard is a static site with no server-side rendering

6. The `--tech` flag

By default, the generated prompt reflects the repository's original technology stack. The `--tech <stack>` flag constrains the prompt to a specific target stack, making it useful for planning a port or rewrite. Multi-word stacks are supported when quoted: single and double quotes group the value into one token and the quote characters are stripped.

How it works

When `--tech` is supplied, the LLM is instructed to:

- Describe how to recreate the repository's *functionality* using the specified technology stack.
- Map the original architecture to equivalent components in the target stack.
- Note where the target stack differs from the original and what adjustments are needed.

Examples

Port a Python project to Rust:

```
/reverse example/flask-api --tech rust
```

Port a Node.js project to Python:

```
/reverse example/express-server --tech python
```

Port a Go project to TypeScript:

```
/reverse example/go-microservice --tech typescript
```

Specify a more detailed stack:

```
/reverse example/legacy-java-app --tech "Rust with axum and SQLx"
```

7. The `--create` flag

The `--create <name>` flag chains the reverse-engineering output into `/spec create`, auto-generating a formal specification from the reverse-engineered prompt.

How it works

1. `/reverse` fetches the repository and generates the synthetic creation prompt.
2. The generated prompt is passed as the feature description to `/spec create <name>`.
3. The spec system creates `specs/<name>/SPEC.md` (EARS requirements), `specs/<name>/PLAN.md` (implementation plan), and `specs/<name>/TESTPLAN.md` (manual test plan).

Examples

Reverse-engineer a repo and create a spec:

```
/reverse thrivethrough/omitme --create omitme-clone
```

Reverse-engineer with a tech constraint and create a spec:

```
/reverse thrivethrough/omitme --tech rust --create omitme-rust-port
```

Reverse-engineer a large project into a spec:

```
/reverse tokio-rs/tokio --create tokio-study
```

What you get

After `--create` finishes, the spec directory contains:

```
specs/omitme-clone/
  SPEC.md          # EARS requirements derived from the creation prompt
  PLAN.md          # Implementation plan with task table
  TESTPLAN.md      # Manual test plan with test cases
```

You can then use the standard `/spec` commands to validate, plan, and implement the spec (see `docs/howtos/spec.md` for details).

8. API interaction

GitHub authentication

A GitHub token is a **prerequisite**: `/reverse` refuses to run without one and reports “No GitHub token configured. Run `/github login` to authenticate, then re-run `/reverse`.” The token is resolved via `/github login` (OAuth device flow, stored in `~/.ragent/github_token`) or the `GITHUB_TOKEN` environment variable. Unauthenticated GitHub API access allows only 60 requests per hour per IP; an authenticated token raises this to 5,000 requests per hour.

```
export GITHUB_TOKEN="ghp_your_token_here"
```

The token does not need any special scopes for reading public repositories. A fine-grained token with read access to public repositories is sufficient.

GitLab authentication

GitLab API calls require a personal access token. The token is resolved in priority order:

1. `GITLAB_TOKEN` environment variable
2. `ragent.json` configuration
3. Encrypted credential database (configured via `/gitlab setup`)

For self-hosted GitLab instances, the host is resolved in priority order: explicit host in the repo identifier (`gitlab:host/namespace/project`), then the `GITLAB_URL` environment variable, then `https://gitlab.com`.

```
export GITLAB_TOKEN="glpat-your_token_here"
export GITLAB_URL="https://gitlab.example.com"
```

Rate limiting

If the API rate limit is exceeded, the command will report an error. Wait for the rate limit window to reset (typically 1 hour) or set `GITHUB_TOKEN` for higher limits.

When using `--depth` for recursive tree fetch, each subdirectory level adds additional API calls (one per subdirectory). Use a token for deeper trees to avoid hitting the rate limit.

Error handling

Error	Cause	What to do
repository not found	Invalid owner/repo or private repo	Check the spelling; set GITHUB_TOKEN for private repos
rate limit exceeded	Too many API requests	Wait for reset or set GITHUB_TOKEN
README not found	Repository has no README	The command continues with metadata and tree only
network error	Connection issue	Check network connectivity and retry
gitlab token missing	No GitLab token configured	Set GITLAB_TOKEN env var or run /gitlab setup
invalid depth	--depth value out of range	Use a value between 1 and 10

Silent degradation on GitHub fetch failures: metadata/tree fetch failures on the GitHub path are swallowed (`unwrap_or_default()`), so an API error produces an empty context fed to the LLM rather than an error message — the `repository not found` / `rate limit exceeded` / `network error` messages only surface on the **GitLab** path. If the generated prompt looks thin, re-run after checking your token and rate limit.

9. CLI equivalents

The TUI `/reverse` command is the primary entry point; the equivalent `ragent reverse` CLI subcommand is not currently registered in the clap CLI (`run`, `serve`, `session`, `memory`, `auth`, `models`, `config`, `research` are the available subcommands). The examples below describe the intended CLI surface and will become accurate when the subcommand ships:

```
# Basic reverse-engineering
```

```
ragent reverse thrivethrough/omitme
```

```
# With technology constraint
```

```
ragent reverse thrivethrough/omitme --tech rust
```

```
# With spec creation
```

```
ragent reverse thrivethrough/omitme --create omitme-clone
```

```
# Full GitHub URL
```

```
ragent reverse https://github.com/thrivethrough/omitme
```

```
# Both flags
```

```
ragent reverse thrivethrough/omitme --tech rust --create omitme-rust-port
```

The CLI command prints the generated prompt to stdout, making it easy to pipe into other tools or save to a file:

```
# Save the generated prompt to a file
```

```
ragent reverse thrivethrough/omitme > omitme-prompt.md
```

```
# Pipe through jq for structured output
```

```
ragent reverse thrivethrough/omitme --tech rust 2>&1 | tee omitme-rust.md
```

10. The `--depth` flag

By default, `/reverse` fetches only the root-level file tree (depth 1). The `--depth <N>` flag controls how many levels of subdirectories are expanded.

Depth	What you get
1 (default)	Root-level files and directories only
2	Root + one level of subdirectories
3	Root + two levels of subdirectories
N (max 10)	Root + N-1 levels of subdirectories

Deeper trees give the LLM more context about the project's module structure, which improves the quality of the generated creation prompt — especially for large repositories with deep nesting. However, each additional level requires more API calls (one per subdirectory), so use a token for higher rate limits when fetching deep trees.

Examples

```
# Default depth (root only)
/reverse thrive/through/omitme

# Two levels deep
/reverse thrive/through/omitme --depth 2

# Maximum depth for a large project
/reverse torvalds/linux --depth 5 --tech rust
```

11. End-to-end examples

Example 1: Understand an unfamiliar project

You find an interesting repository and want to understand what it does:

```
/reverse thrive/through/omitme
```

The generated prompt gives you a comprehensive overview of the project's purpose, architecture, and technology stack — without needing to browse the code yourself.

Example 2: Plan a port to a different language

You want to port a Python project to Rust:

```
/reverse example/flask-api --tech rust
```

The generated prompt describes how to recreate the Flask API's functionality using Rust, mapping each component to its Rust equivalent (e.g. Flask → Axum, SQLAlchemy → SQLx, Celery → Tokio tasks).

Example 3: Create a spec from a reference implementation

You want to create a formal spec based on an existing open-source project:

```
/reverse tokio-rs/tokio --create tokio-study
```

This generates a spec under `specs/tokio-study/` with EARS requirements, an implementation plan, and a test plan — all derived from the reverse-engineered creation prompt.

Example 4: Study a well-architected project

You want to study how a well-known project is structured:

```
/reverse burntsushi/regex
```

The generated prompt breaks down the regex crate's architecture, module structure, and key design decisions, giving you a learning roadmap.

Example 5: Plan a rewrite with a specific stack

You want to rewrite a legacy application using a modern stack:

```
/reverse example/legacy-monolith --tech "Rust with axum, SQLx, and Redis"
```

The quotes are required here: an unquoted value would stop at the first space. The generated prompt maps the legacy monolith's functionality to a modern Rust microservice architecture with specific crate recommendations.

Example 6: Reverse-engineer and immediately implement

Combine `/reverse` with `/spec impl` for a full pipeline:

```
# Step 1: Reverse-engineer and create a spec
```

```
/reverse thrive-through/omitme --tech rust --create omitme-rust
```

```
# Step 2: Validate the generated spec
```

```
/spec validate omitme-rust
```

```
# Step 3: Preview the implementation plan
```

```
/spec impl omitme-rust --dry-run
```

```
# Step 4: Implement
```

```
/spec impl omitme-rust
```

Example 7: Compare two implementations

Reverse-engineer two competing projects to compare their approaches:

```
/reverse project-a/repo --create study-a
```

```
/reverse project-b/repo --create study-b
```

Then compare the generated specs to understand the architectural differences between the two projects.

Example 8: Generate a prompt for a specific use case

You want a creation prompt focused on the testing infrastructure of a project:

```
/reverse example/well-tested-app --tech rust-proptest
```

The generated prompt emphasises how to recreate the project's testing approach using the specified testing tools.

Example 9: Reverse-engineer a GitLab repository

```
# Self-hosted GitLab with nested namespace
```

```
/reverse gitlab:gitlab.example.com/group/subgroup/project --depth 3
```

```
# GitLab.com with spec creation
```

```
/reverse gitlab:my-namespace/my-project --create my-project-clone
```

```
# Full GitLab URL
```

```
/reverse https://gitlab.com/my-namespace/my-project --tech python --create my-py-port
```

12. Tips for good results

- **Use specific tech stacks.** `--tech rust` is good, but unquoted values end at the first space — quote multi-word stacks (`--tech "Rust with axum and SQLx"`) or use hyphenated descriptors like `--tech rust-axum-sqlx` for more targeted prompts.
 - **Review the generated prompt before acting.** The synthetic prompt is a starting point — review it for accuracy and adjust before feeding it to a coding agent.
 - **Use `--create` for structured workflows.** Chaining into `/spec create` gives you a formal spec, plan, and test plan — much more actionable than a raw prompt.
 - **Set `GITHUB_TOKEN` for frequent use.** A token is required to run the command at all; the unauthenticated API allows only 60 requests/hour, authenticated is 5,000/hour.
 - **Use full URLs for clarity.** `/reverse https://github.com/owner/repo` is unambiguous; `/reverse owner/repo` is faster to type.
 - **Combine with `/spec` commands.** After `--create`, use `/spec validate`, `/spec plan`, `/spec tasks`, and `/spec impl` for a complete spec-driven workflow.
 - **Use `--tech` for port planning.** When planning a port, always specify the target stack so the prompt maps to the right technologies.
 - **Try different models.** The generated prompt quality depends on the selected LLM model. Try different models (e.g. Claude, GPT-4) to compare output quality.
 - **Study well-architected projects.** Use `/reverse` on projects you admire to learn about their architecture and design patterns.
-

13. Troubleshooting

Symptom	Likely cause	What to do
repository not found	Invalid owner/repo or private repo	Check spelling; set GITHUB_TOKEN for private repos
rate limit exceeded	Too many unauthenticated API calls	Set GITHUB_TOKEN environment variable
README not found	Repository has no README	The command continues with metadata and tree only
Empty or short prompt	LLM model not configured or down	Check provider setup with /models
network error	Connection issue	Check network connectivity and retry
invalid repository format	Malformed input	Use owner/repo or a full GitHub URL
spec already exists	--create target name already used	Choose a different name or delete the existing spec
Prompt mentions wrong technologies	--tech not specified or too vague	Use a more specific --tech value

14. Workflow integration

The /reverse command integrates with ragent's spec management system to provide a complete reverse-engineering-to-implementation pipeline:

```
/reverse <owner/repo>
```

```
GitHub API fetch
(metadata + tree
+ README)
```

```
LLM generates
creation prompt
```

```
--create <name>
Display prompt          /spec create
in chat window          generates SPEC.md,
                        PLAN.md, TESTPLAN.md
```

```
/spec validate
/spec plan
/spec impl
```

Full pipeline example

```
# 1. Reverse-engineer a project
/reverse thrive/omitme --tech rust --create omitme-rust

# 2. Validate the generated spec
/spec validate omitme-rust

# 3. Generate task list
/spec tasks omitme-rust

# 4. Preview implementation
/spec impl omitme-rust --dry-run

# 5. Implement
/spec impl omitme-rust
```

This pipeline takes you from an unfamiliar GitHub repository to a running implementation in five commands.

End of manual.