/bench

Benchmark runner: /bench list|init |show|run |status|open last|cancel

Overview

The `/bench` command drives the ragent benchmark harness (`ragent-bench` crate). The TUI arm parses its arguments with `ragent_bench::parse_bench_command` and dispatches the parsed command to the bench runtime; everything below the `/bench` prefix is handled by the bench crate, not the TUI itself.

The full guide (concepts, every subcommand, run flags, suites, workbook layout, workflows and troubleshooting) is docs/userdocs/bench.md.

Syntax

```
/bench list
/bench init <suite-or-all-or-full> [--full] [--language LANG] [--force-download] [--verify-only]
/bench show
/bench run <suite-or-profile-or-all> [flags]
/bench status
/bench open last
/bench cancel
/bench help
```

- Subcommands: Command reference.
- `run` flags (`--limit`/`--cap`, `--samples`, `--subset`, `--release`, `--scenario`, `--language`, `--temperature`, `--top-p`, `--max-tokens`, `--deterministic`, `--since`, `--until`, `--resume`, `--no-exec`, `--yes`): Run flags.
- `help` prints the bench usage help.

Cancellation

`/bench cancel` prints a confirmation line and the status bar shows `[wait] bench: cancellation requested` while the flag is set. When no run is active, `cancel` prints “No active benchmark run.” instead.

Related

- `/model` - verify the provider and model a bench run will target
- `/telemetry` - runtime metric counters recorded during runs
- `cargo bench -p ragent-tui` and friends for criterion micro-benchmarks