

/bench

Benchmark runner: `/bench list|init |show|run |status|open last|cancel`

Overview

The `/bench` command drives the `ragent-benchmark` harness (`ragent-bench` crate). It initialises benchmark suites, runs them against configured LLM models, and inspects results. The TUI arm parses its arguments with `ragent_bench::parse_bench_command` and dispatches the parsed command to the bench runtime; everything below the `/bench` prefix is handled by the bench crate, not the TUI itself.

Suites measure tool-calling behaviour, code quality, and end-to-end agent runs. Results are stored under the bench output directory so later run invocations can resume or compare against previous releases.

Syntax

```
/bench list
/bench init <suite-or-all-or-full> [--full] [--language LANG] [--force-download] [--verify-only]
/bench show <suite>
/bench run <suite-or-profile-or-all> [flags]
/bench status
/bench open last
/bench cancel
/bench help
```

Options / Subcommands

Form	Description
<code>list</code>	List the registered benchmark suites and profiles
<code>init</code>	Initialise fixtures for a suite, all suites, or the full corpus
<code><suite-or-all-or-full></code>	
<code>init ... --full</code>	Build the complete (large) fixture set instead of the default sample
<code>init ... --language LANG</code>	Restrict fixture generation to one language
<code>init ...</code>	Re-download fixture inputs even if already cached
<code>--force-download</code>	
<code>init ... --verify-only</code>	Verify existing fixtures without generating or downloading
<code>run</code>	Execute a suite, a named profile, or every suite
<code><suite-or-profile-or-all></code>	
<code>run ... --limit N / --cap N</code>	Cap the number of benchmark cases executed
<code>run ... --samples K</code>	Number of samples per case
<code>run ... --subset NAME</code>	Run only the named subset
<code>run ... --release</code>	Tag the run against a release version
<code>VERSION</code>	
<code>run ... --scenario NAME</code>	Run one scenario
<code>run ... --language LANG</code>	Restrict the run to one language
<code>run ... --temperature F</code>	Override sampling temperature
<code>run ... --top-p F</code>	Override top-p sampling

Form	Description
<code>run ... --max-tokens N</code>	Override the completion token budget
<code>run ... --deterministic</code>	Force deterministic sampling settings
<code>run ... --since</code> YYYY-MM-DD	Only include data on or after this date
<code>run ... --until</code> YYYY-MM-DD	Only include data on or before this date
<code>run ... --resume</code>	Resume a previously interrupted run
<code>run ... --no-exec</code>	Parse and plan the run without executing tool calls
<code>run ... --yes</code>	Auto-approve prompts during the run
<code>status</code>	Show whether a bench run is active and its last known state
<code>open last</code>	Open the most recent bench result report
<code>cancel</code>	Request cancellation of the active benchmark run
<code>help</code>	Print the bench usage help

Examples

From: `/bench list`

Lists every benchmark suite with a short description and case count.

From: `/bench init all`

Initialises fixtures for every suite. Progress lines report each suite as it is generated; use `--verify-only` afterwards to re-check them.

From: `/bench run codebench --language rust --limit 10`

Runs ten rust cases from the codebench suite with the current provider and model.

From: `/bench run all --deterministic --max-tokens 2048 --release 1.0.95`

Runs every suite with deterministic sampling, a 2048-token cap, and stores the result under release 1.0.95 for later comparison.

From: `/bench status`

Reports the state of any in-flight run; prints the last run summary when idle.

From: `/bench cancel`

Cancellation requested for the active benchmark run.

When no run is active, cancel prints “No active benchmark run.” instead.

Output

`init` and `run` stream progress into the message window as the harness works: suite names, case indices, pass/fail outcomes, and aggregate scores. `status` prints the run state; `open last` shows the stored result report; `cancel` prints a confirmation line and the status bar shows `[wait] bench: cancellation requested` while the flag is set.

Related

- `/models` - verify the provider and model a bench run will target
- `/telemetry` - runtime metric counters recorded during runs

- `cargo bench -p ragent-tui` and friends for criterion micro-benchmarks
- `docs/performance/benchmark-guide.md`