

/cost

Show session token usage and estimated cost

Overview

`/cost` summarises the LLM traffic recorded for the current session: number of completed request samples, session duration, total input/output tokens, and an estimated cost in USD. Cost is computed per sample from the resolved model's configured per-million-token input and output rates; samples whose model cannot be resolved contribute tokens but zero cost. The report also breaks totals down by provider, sorted alphabetically.

If no completed LLM responses have been recorded yet, the command says so instead of printing an empty summary.

Syntax

`/cost`

Options / Subcommands

Form	Description
<code>/cost</code>	Show session token usage and estimated cost

No arguments or subcommands.

Examples

`/cost`

Typical report:

From: `/cost`

Samples: 14

Session duration: 12m 31s

Total tokens: 48210 input / 3918 output

Estimated cost: \$0.412370

By provider:

- anthropic: \$0.412370 (48210 in / 3918 out, 14 samples)
- ollama: \$0.000000 (2150 in / 260 out, 2 samples)

`/cost`

Nothing recorded yet:

No completed LLM responses yet for this session.

Output

- Assistant bubble From: `/cost` with: `Samples:`, `Session duration:`, `Total tokens:` `N` input / `N` output, `Estimated cost:` `$X.XXXXXX`, and an optional `By provider:` block listing per-provider cost, tokens, and sample counts.
- Session duration is derived from the stored session's creation time (hours/minutes/seconds format).
- No-data case: assistant bubble `No completed LLM responses yet for this session.` and status `cost unavailable.`
- Normal status bar: `cost summary.`

Related

- `/about` - application version info
- `/startup` - startup timing report
- Model costs are configured per model in `ragent.json`