

/thinking

Switch the current thinking level: `/thinking auto|off|low|medium|high`

Overview

`/thinking` reports or changes the reasoning (extended thinking) level of the currently selected model. Bare `/thinking` prints the current level and the levels the model supports; passing a level switches it and persists the choice. The command is model-gated: it refuses to run before a model has been selected with `/model`, and it validates the requested level against the levels the active model actually supports.

Syntax

```
/thinking          Show current level and supported levels
/thinking <level>  Set the thinking level
```

Accepted `<level>` values: `auto`, `off`, `low`, `medium`, `high`.

Options / Subcommands

Form	Description
<code>/thinking</code>	Print the current thinking level and the supported levels
<code>/thinking auto</code>	Let the provider default decide thinking behaviour
<code>/thinking off</code>	Disable extended thinking
<code>/thinking low</code>	Low reasoning effort
<code>/thinking medium</code>	Medium reasoning effort
<code>/thinking high</code>	High reasoning effort

Behaviour

- **Model required.** With no model selected the command prints `[warn] No model selected - use /model to choose and does nothing.`
- **Validation.** The requested level is validated against the active model's supported level set. An unsupported level prints `[warn] Thinking level '<level>' is not supported by the active model.`
- **Thinking-incapable models.** If the active model reports no configurable thinking levels, only `off` is accepted; any other level prints `[warn] Active model does not support configurable thinking.`
- **Persistence.** The chosen level is stored against the selected model and is used for subsequent turns in the session.
- An invalid token (not one of the five levels) prints `Usage: /thinking [auto|off|low|medium|high].`

Examples

```
/thinking
```

Prints the current and supported levels, for example: `Current: 'medium' and Supported: 'auto, low, medium, high'.`

`/thinking high`

Switches the model to high reasoning effort; status line becomes `thinking: high`.

`/thinking off`

Turns extended thinking off for the active model.

`/thinking turbo`

Invalid level - prints the usage line in the status bar.

Output

- Bare form: chat text with the current level and the supported level list, status line `thinking`.
- Set form: no chat text; status line `thinking: <level>`.
- Refusals: status line only (`[warn]` No model selected - use `/model` to choose, unsupported-level warning, unsupported-model warning, or the usage line).

Related

- `/model` - select the model whose thinking level you configure
- `/provider` - change provider configuration
- Provider `thinking` section in `ragent.json` - per-provider and per-model thinking defaults (`enabled`, `level`, `budget_tokens`)