

LLM Provider Comparison — OpenCode vs Codex CLI vs Claude Code vs GitHub Copilot vs ragent

Last updated: 2025-07

This document compares the LLM provider support across five AI coding agents for the terminal.

Summary

Feature	OpenCode	Codex CLI	Claude Code	GitHub Copilot CLI	ragent
Total providers	75+	1 (OpenAI)	1 (Anthropic)	3	10
Extensible	✓ (AI SDK)	✗	✗	✓ (BYOK, plugins)	✓ (Generic OpenAI)
Local models	✓ (Ollama, llama.cpp, LM Studio)	✓ (via OpenRouter/Ollama proxy)	✗	✓ (local model support)	✓ (Ollama)
Cloud-only models	✓	✓	✓	✓	✓
Custom/OpenAI-compatible	✓ (baseUrl override)	✓ (via proxy)	✗	✓ (BYOK)	✓ (Generic OpenAI)
Open source	✓	✓	✗	✗	✓

Provider Matrix

Provider	OpenCode	Codex CLI	Claude Code	GitHub Copilot CLI	ragent
Anthropic	✓	✗	✓ (native)	✓	✓
OpenAI	✓	✓ (native)	✗	✓	✓
Google Gemini	✓	🔒🔒🔒	✗	✓	✓
Ollama (local)	✓	✗*	✗	✓	✓
Ollama Cloud	✓	✗	✗	✗	✓

Provider	OpenCode	Codex CLI	Claude Code	GitHub Copilot CLI	ragent
Hugging Face	✓	✗	✗	✗	✓
GitHub Copilot	✓	✗	✗	✓ (native)	✓
Generic OpenAI-compatible	✓ (baseUrl)	✓ (baseUrl)	✗	✓ (BYOK)	✓ (native)
Azure AI Foundry	✗	✗	✗	✓	✓
Azure Resource (File)	✗	✗	✗	✗	✓
Amazon Bedrock	✓	✗	✓	✗	✗
Google Vertex AI	✗	✗	✓	✗	✗
Azure OpenAI	✓	✗	✗	✓	✗
DeepSeek	✓	✗	✗	✗	✗
Groq	✓	✗	✗	✗	✗
Cerebras	✓	✗	✗	✗	✗
Together AI	✓	✗	✗	✗	✗
Fireworks AI	✓	✗	✗	✗	✗
NVIDIA	✓	✗	✗	✗	✗
xAI (Grok)	✓	✗	✗	✗	✗
OpenRouter	✓	✗	✗	✗	✗
llama.cpp	✓	✗	✗	✗	✗
LM Studio	✓	✗	✗	✗	✗
Cloudflare Workers AI	✓	✗	✗	✗	✗
Scaleway	✓	✗	✗	✗	✗
DigitalOcean	✓	✗	✗	✗	✗
Moonshot AI	✓	✗	✗	✗	✗
MiniMax	✓	✗	✗	✗	✗
302.AI	✓	✗	✗	✗	✗
SAP AI Core	✓	✗	✗	✗	✗
Vercel AI Gateway	✓	✗	✗	✗	✗

Provider	OpenCode	Codex CLI	Claude Code	GitHub Copilot CLI	ragent
Venice AI	✓	✗	✗	✗	✗
IO.NET	✓	✗	✗	✗	✗
GitLab Duo	✓	✗	✗	✗	✗
FrogBot	✓	✗	✗	✗	✗
Baseten	✓	✗	✗	✗	✗
Helicone	✓	✗	✗	✗	✗
Cortecs	✓	✗	✗	✗	✗
OVHcloud AI Endpoints	✓	✗	✗	✗	✗
STACKIT	✓	✗	✗	✗	✗
Nebius Token Factory	✓	✗	✗	✗	✗
Z.AI	✓	✗	✗	✗	✗

* Codex CLI can work with Ollama via a local OpenAI-compatible proxy, but has no native Ollama integration.

Detailed Agent Breakdown

OpenCode

- **Total providers:** 75+
- **Architecture:** Uses the [AI SDK](#) and [Models.dev](#) for broad provider coverage
- **Local model support:** Ollama, llama.cpp, LM Studio
- **Custom providers:** Any OpenAI-compatible endpoint via [baseUrl](#) override
- **Proprietary offerings:** OpenCode Zen (curated models) and OpenCode Go (low-cost subscription)
- **Configuration:** [opencode.json](#) with provider-specific sections; API keys stored in [~/.local/share/opencode/auth.json](#) via [/connect](#) command
- **Notable:** Widest provider coverage of all five agents; supports niche and regional providers (302.AI, STACKIT, OVHcloud, Scaleway, etc.)

OpenAI Codex CLI

- **Total providers:** 1 (OpenAI native)
- **Architecture:** Direct OpenAI API integration
- **Supported models:** GPT-5.4 (default), GPT-5.4 Mini, GPT-5.3-Codex, GPT-5.2 Codex, GPT-5.2, o4-mini, Codex Mini

- **Local model support:** None native; works via OpenAI-compatible proxies (Ollama, LM Studio, OpenRouter)
- **Custom providers:** Can point to any OpenAI-compatible `baseURL`
- **Configuration:** `--provider` flag and `CODEX_HOME/config.json`
- **Pricing:** Free open-source tool; model costs through OpenAI API or ChatGPT subscription tiers (Free → Pro \$200/mo)
- **Notable:** Sandboxed execution environment; tightly integrated with OpenAI model ecosystem; limited multi-provider support by design

Claude Code

- **Total providers:** 1 (Anthropic native) + 2 cloud routes
- **Architecture:** Direct Anthropic API integration
- **Supported models:** Claude Opus 4, Claude Sonnet 4, Claude Haiku 3.5
- **Cloud routes:** Amazon Bedrock, Google Vertex AI (via `CLAUDE_CODE_USE_BEDROCK` / `CLAUDE_CODE_USE_VERTEX` env vars)
- **Local model support:** None
- **Custom providers:** None native; third-party proxies (Portkey, OpenRouter) can bridge to other providers
- **Configuration:** Environment variables for provider routing
- **Pricing:** Anthropic API pricing or Claude subscription (Pro \$20/mo, Max \$100/mo)
- **Notable:** Best-in-class reasoning with Claude Opus; closed source; strong single-provider focus; third-party workarounds exist for multi-provider use

GitHub Copilot CLI

- **Total providers:** 3 (OpenAI, Anthropic, Google) + BYOK
- **Architecture:** GitHub-hosted model routing with enterprise controls
- **Supported models:** GPT-4o, GPT-4.1, Claude 3.5 Sonnet, Gemini 2.5 Pro, Gemini 2.5 Flash, o3-mini, o4-mini, and auto-selection mode
- **BYOK support:** Bring-your-own-key for external providers via plugins
- **Local model support:** Yes, with local LLM endpoints
- **Custom providers:** Plugin system and BYOK for custom endpoints
- **Configuration:** GitHub authentication + settings; enterprise policies for model governance
- **Pricing:** Copilot Free / Pro (\$19/mo) / Business (\$39/mo) / Enterprise (\$39/user/mo); premium models have request multiplier costs
- **Notable:** Enterprise-ready with FedRAMP support, model governance policies, and usage metering; plugin/extension ecosystem for customisation; closed source

ragent

- **Total providers:** 10 (8 cloud + 2 local)
- **Architecture:** Custom `Provider` trait with `ProviderRegistry`; each provider is a dedicated Rust module
- **Native providers:**

1. **Anthropic** — Claude models with thinking/reasoning support
 2. **OpenAI** — GPT models
 3. **Google Gemini** — Gemini models
 4. **Ollama** — Local model inference (any GGUF model)
 5. **Ollama Cloud** — Cloud-hosted Ollama service
 6. **Hugging Face** — HF Inference API models
 7. **GitHub Copilot** — Copilot models (with request multiplier support)
 8. **Generic OpenAI** — Any OpenAI-compatible endpoint (custom `base_url` and API key)
 9. **Azure AI Foundry** — Azure-hosted models with dynamic model discovery
 10. **Azure Resource (File)** — Azure file/resource provider with `azureresources.json` support
- **Local model support:** Ollama (native, first-class)
 - **Custom providers:** Generic OpenAI-compatible endpoint with configurable `base_url` — covers LM Studio, vLLM, LocalAI, and any OpenAI API-compatible server
 - **Configuration:** `ragent.json` (or `.jsonc`) with per-provider sections; environment variable support (`ANTHROPIC_API_KEY`, `OPENAI_API_KEY`, etc.)
 - **Pricing:** Free open-source tool; model costs through individual provider APIs
 - **Notable:** Rust-native performance; thinking/reasoning budget control per model; unique Azure AI Foundry and Azure Resource providers; compiled as single statically-linked binary with zero runtime dependencies

Feature Comparison

Feature	OpenCode	Codex CLI	Claude Code	Copilot CLI	ragent
Open source	✓	✓	✗	✗	✓
Single binary	✗ (Go)	✗ (Node)	✗ (Node)	✗ (Node)	✓ (Rust)
Zero runtime deps	✗	✗	✗	✗	✓
Thinking/reasoning config	✗	✗	✗	✗	✓
MCP client	✓	✗	✓	✓	✓
Local Ollama (native)	✓	✗	✗	✓	✓
Azure AI Foundry	✗	✗	✗	✓	✓
Amazon Bedrock	✓	✗	✓	✗	✗
Google Vertex AI	✗	✗	✓	✗	✗
Generic OpenAI endpoint	✓	✓	✗	✓	✓
Hugging Face	✓	✗	✗	✗	✓
Sandbox execution	✗	✓	✗	✗	✗
Team/swarm agents	✗	✗	✗	✓	✓

Feature	OpenCode	Codex CLI	Claude Code	Copilot CLI	ragent
Codebase indexing	✗	✗	✗	✓	✓

Key Takeaways

1. **Broadest coverage:** OpenCode leads with 75+ providers via the AI SDK, making it ideal for teams that need niche or regional cloud providers.
2. **Deepest single-provider integration:** Claude Code offers the best experience for Anthropic models with native Bedrock and Vertex AI routing, but is locked to a single model family.
3. **Enterprise governance:** GitHub Copilot CLI provides the strongest enterprise story with FedRAMP support, model governance policies, and usage metering across its three supported providers.
4. **Balanced multi-provider with local-first:** ragent offers a curated set of 10 providers covering all major cloud providers plus native Ollama support, with the unique advantage of being a single static binary and offering thinking/reasoning budget controls per model.
5. **OpenAI ecosystem focus:** Codex CLI is purpose-built for OpenAI models with a sandboxed execution model, making it the best choice for teams standardised on OpenAI.
6. **Local model support:** OpenCode, Copilot CLI, and ragent all support local inference via Ollama. Codex CLI can work through proxies. Claude Code has no local model path.